

# CSE3081 Design and Analysis of Algorithms

Dept. of Computer Engineering,  
Sogang University

This material contains text and figures from other lecture slides. Do not post it on the Internet.

# Chapter 8. Sorting in Linear Time

---

# Sorting Algorithms

- Until now, we have seen sorting algorithms with time complexity  $\Theta(n^2)$  and  $\Theta(n \log n)$ .

Algorithm	Worst-case	Average-case
Insertion Sort	$\Theta(n^2)$	$\Theta(n^2)$
Selection Sort	$\Theta(n^2)$	$\Theta(n^2)$
Bubble Sort	$\Theta(n^2)$	$\Theta(n^2)$
Merge Sort	$\Theta(n \lg n)$	$\Theta(n \lg n)$
Heapsort	$\Theta(n \lg n)$	$\Theta(n \lg n)$
Quicksort	$\Theta(n^2)$	$\Theta(n \lg n)$

- These algorithms have one thing in common: sorting is done based only on **comparisons between the input elements**.
- We call these algorithms **comparison sorts** (comparison-based sorting).

## 8.1 Lower Bounds for Sorting

---

# Comparison Sorts


- A comparison sort uses only comparisons between elements to gain order information about an input sequence  $\langle a_1, a_2, \dots, a_n \rangle$ .
- Given two elements  $a_i$  and  $a_j$ , it performs tests such as  $a_i < a_j$ ,  $a_i \leq a_j$ ,  $a_i = a_j$ ,  $a_i \geq a_j$ ,  $a_i > a_j$  in order to determine their relative order.
- It may not inspect the values of the elements or gain order information about them in any other way.

# Comparison Sorts: Assumptions

- Without loss of generality, we assume that all input elements are distinct.
  - The same proof will apply for the case when elements are not distinct.
- With the above assumption, we can assume that all comparisons have the form  $a_i \leq a_j$ .
  - $a_i \leq a_j$ ,  $a_i \geq a_j$ ,  $a_i > a_j$ , and  $a_i < a_j$  all yield identical information about the relative order of  $a_i$  and  $a_j$ .


# Comparison Sorts: Decision-Tree Model

- Decision Tree
  - A full binary tree that represents the comparisons between elements that are performed by a particular sorting algorithm
  - **Full binary tree**: a tree in which each node is either a leaf or has both children



# Comparison Sorts: Decision-Tree Model


- Each internal node is annotated by  $i:j$  for some  $i$  and  $j$  in the range  $1 \leq i, j \leq n$ , where  $n$  is the number of elements in the input sequence.
- Each leaf node is annotated by a permutation  $\langle \pi(1), \pi(2), \dots, \pi(n) \rangle$ .
- Indices in the internal nodes and the leaves refer to the original positions of the array elements at the start of the sorting algorithm.
- The execution of the comparison sort corresponds to tracing a simple path from the root of the decision tree down to a leaf.
- Each of the  $n!$  permutations on  $n$  elements must appear as at least one of the leaves of the decision tree.





# Comparison Sorts: Lower Bound

- The worst-case number of comparisons
  - The length of the longest simple path from the root of a decision tree to any of its reachable leaves
  - The height of its decision tree



# Comparison Sorts: Lower Bound

- Theorem 8.1
    - Any comparison sort algorithm requires  $\Omega(n \log n)$  comparisons in the worst case.
  - Proof
    - Consider a decision tree of height  $h$  with  $l$  reachable leaves corresponding to a comparison sort with  $n$  elements. Thus, we have  $n! \leq l$ .
    - Since a binary tree of height  $h$  has no more than  $2^h$  leaves, we have
    - $n! \leq l \leq 2^h$
    - Taking logarithms,
    - $h \geq \log n! = \Omega(n \log n)$
- $$n! = \sqrt{2\pi n} \left(\frac{n}{e}\right)^n e^{\alpha_n}$$

Stirling's approximation
- Heapsort and merge sort are asymptotically best strategies for comparison sort!

## 8.2 Counting Sort

---

# Counting Sort

- Assumption: Each of the  $n$  input elements is an integer in the range 0 to  $k$ , for some integer  $k$ .
- Under this assumption, the counting sort runs in  $\Theta(n + k)$  time.
  - If  $k = O(n)$ , counting sort runs in  $\Theta(n)$  time.

# Counting Sort: Strategy

- For each input element, counting sort first determines the number of elements less than or equal to  $x$ .
- It then uses this information to place element  $x$  directly into its position in the output array.
  - If 17 elements are less than or equal to  $x$ , then  $x$  belongs in output position 17.
  - We need slight modification to handle elements with the same value, so that they are not placed in the same position.







# Counting Sort: Algorithm

- Algorithm COUNTING-SORT

```
COUNTING-SORT( $A, n, k$ )
1  let  $B[1:n]$  and  $C[0:k]$  be new arrays
2  for  $i = 0$  to  $k$ 
3       $C[i] = 0$ 
4  for  $j = 1$  to  $n$ 
5       $C[A[j]] = C[A[j]] + 1$ 
6  //  $C[i]$  now contains the number of elements equal to  $i$ .
7  for  $i = 1$  to  $k$ 
8       $C[i] = C[i] + C[i - 1]$ 
9  //  $C[i]$  now contains the number of elements less than or equal to  $i$ .
10 // Copy  $A$  to  $B$ , starting from the end of  $A$ .
11 for  $j = n$  downto 1
12      $B[C[A[j]]] = A[j]$ 
13      $C[A[j]] = C[A[j]] - 1$  // to handle duplicate values
14 return  $B$ 
```

# Counting Sort: Operation Example

- Algorithm COUNTING-SORT



# Counting Sort: Analysis and Running Time

- Algorithm COUNTING-SORT
  - Lines 2-3: initializes array  $C$  to all zeros.
  - Lines 4-5: makes a pass over the array  $A$  to inspect each input element. Each time it finds an input element whose value is  $i$ , it increments  $C[i]$ .
  - Lines 7-8: determines for each  $i = 0, 1, \dots, k$  how many input elements are less than or equal to  $i$  by keeping a running sum of the array  $C$ .
  - Lines 11-13: makes another pass over  $A$ , but in reverse, to place each element  $A[j]$  into its correct sorted position in the output array  $B$ .
    - Line 13 is needed to handle duplicate values
- Running time of COUNTING-SORT
  - Lines 2-3:  $\Theta(k)$
  - Lines 4-5:  $\Theta(n)$
  - Lines 7-8:  $\Theta(k)$
  - Lines 11-13:  $\Theta(n)$
  - Overall:  $\Theta(k + n)$



# Comparison Sort: Remarks


- In practice, we can use counting sort when we have  $k = O(n)$ , in which case the running time is  $\Theta(n)$ .
- Counting sort can be asymptotically faster than  $\Theta(n \log n)$  because it is **not a comparison sort**.
  - Counting sort uses actual values of the elements instead of comparisons.
- Counting sort is **stable**
  - Elements with the same value appear in the output array in the same order as they do in the input array.
  - It breaks ties between two elements by the rule that whichever element appears first in the input array appears first in the output array.
  - Property of stability is important only when satellite data are carried around with the element being sorted.
  - Counting sort's stability is important for another reason: it is often used as a subroutine in radix sort. For radix sort to work correctly, counting sort must be stable.

## 8.3 Radix Sort

---

# Radix Sort

- Used by the card-sorting machines for punch cards.
  - The cards have 80 columns, and in each column a machine can punch a hole in one of 12 places.
  - The sorter can be mechanically programmed to examine a given column of each card in a deck and distribute the card into one of 12 bins depending on which place has been punched.
  - An operator can then gather the cards bin by bin, so that cards with the first place punched are on top of cards with the second place punched, and so on.




# Radix Sort: Strategy

- The radix sort solves the problem of card sorting by sorting on the **least significant digit** first.
- The algorithm then combines the cards into a single deck, with the cards in the 0 bin preceding the cards in the 1 bin preceding the cards in the 2 bin, and so on.
- Then it sorts the entire deck again on the second-least significant digit and recombines the deck in a like manner.
- This process continues until the cards have been sorted on all  $d$  digits.
- Remarkably, at that point the cards are fully sorted on the  $d$ -digit number.
- Thus, only  **$d$  passes through the deck are required** to sort.

# Radix Sort: Example and Requirements

- Radix sort



- In order for radix sort to work correctly, the digit sorts must be **stable**.
  - Suppose 457 precedes 458 after sorting the least significant digit.
  - When sorting the second least significant digit, their order must stay the same.

# Radix Sort: Algorithm

- Algorithm RADIX-SORT

```
RADIX-SORT( $A, n, d$ )
```

```
1  for  $i = 1$  to  $d$ 
```

```
2      use a stable sort to sort array  $A[1 : n]$  on digit  $i$ 
```

- For stable sort, COUNTING-SORT is commonly used.

# Radix Sort: Running Time

- Lemma 8.3
  - Given  $n$   $d$ -digit numbers in which each digit can take on up to  $k$  possible values, RADIX-SORT correctly sorts these numbers in  $\Theta(d(n + k))$  time if the stable sort it uses takes  $\Theta(n + k)$  time.
- Proof
  - Each pass over  $n$   $d$ -digit numbers takes  $\Theta(n + k)$  time.
  - There are  $d$  passes.
- When  $d$  is constant and  $k = O(n)$ , we can make radix sort run in linear time.

# Radix Sort: Breaking a Number into Digits

- Lemma 8.4
  - Given  $n$   $b$ -bit numbers and any positive integer  $r \leq b$ , RADIX-SORT correctly sorts these numbers in  $\Theta\left(\left(\frac{b}{r}\right)(n + 2^r)\right)$  time if the stable sort it uses takes  $\Theta(n + k)$  time for inputs in the range 0 to  $k$ .
- Proof
  - For a value  $r \leq b$ , view each key as having  $d = \lceil b/r \rceil$  digits of  $r$  bits each.
  - Each digit is an integer in the range 0 to  $2^r - 1$ , so that we can use counting sort with  $k = 2^r - 1$ .
  - (For example, a 32-bit word can be viewed as having four 8-bit digits, so that  $b = 32$ ,  $r = 8$ ,  $k = 2^r - 1 = 255$ , and  $d = \frac{b}{r} = 4$ .)
  - Each pass of counting sort takes  $\Theta(n + k) = \Theta(n + 2^r)$  time and there are  $d$  passes, for a total running time of  $\Theta(d(n + 2^r)) = \Theta\left(\left(\frac{b}{r}\right)(n + 2^r)\right)$ .



# Radix Sort: Breaking a Number into Digits

- Given  $n$  and  $b$ , what value of  $r \leq b$  minimizes the expression  $\left(\frac{b}{r}\right)(n + 2^r)$ ?
  - As  $r$  decreases, the factor  $b/r$  increases, but as  $r$  increases so does  $2^r$ .
  - The answer depends on whether  $b < \lfloor \log n \rfloor$ .
  - If  $b < \lfloor \log n \rfloor$ , then  $r \leq b$  implies  $(n + 2^r) = \Theta(n)$ .
  - Thus, choosing  $r = b$  yields a running time of  $\left(\frac{b}{b}\right)(n + 2^b) = \Theta(n)$ , which is optimal.
  - If  $b \geq \lfloor \log n \rfloor$ , then choosing  $r = \lfloor \log n \rfloor$  gives the best running time.
  - Choosing  $r = \lfloor \log n \rfloor$  yields a running time of  $\Theta\left(\frac{bn}{\log n}\right)$ .
  - As  $r$  increases above  $\lfloor \log n \rfloor$ , the  $2^r$  term in the numerator increases faster than the  $r$  term in the denominator, and so increasing  $r$  above  $\lfloor \log n \rfloor$  yields a running time of  $\Omega\left(\frac{bn}{\log n}\right)$ .
  - If instead  $r$  were to decrease below  $\lfloor \log n \rfloor$ , then the  $b/r$  term increases and the  $n + 2^r$  term remains at  $\Theta(n)$ .

# Radix Sort: A Good Choice?

- Is radix sort preferable to a comparison-based sorting algorithm, such as quicksort?
- If  $b = O(\log n)$ , and  $r \approx \log n$ , then radix sort's running time is  $\Theta(n)$ , which appears to be better than quicksort's expected running time of  $\theta(n \log n)$ .
- However, the constant factors hidden in the  $\Theta$ -notation differ.
- Although radix sort may make fewer passes than quicksort over the  $n$  keys, each pass of radix sort may take significantly longer.
- Thus, which sorting algorithm to prefer depends on the input data and characteristics of the implementations and underlying machine
  - quicksort often uses hardware caches more effectively than radix sort
  - radix sort that uses counting sort does not sort in place, which many of the  $\Theta(n \log n)$ -time comparison sorts do. Thus, when primary memory storage is at a premium, an in-place algorithm such as quicksort could be the better choice.

## 8.4 Bucket Sort

---

# Bucket Sort: Introduction

- Bucket sort is a sorting algorithm that has an average-case running time of  $O(n)$ .
- Bucket sort assumes that the input is drawn from **a uniform distribution**.
  - Bucket sort is fast because of the assumption.
- Assumptions: Counting sort vs. Bucket sort
  - Counting sort: input consists of integers in a small range
  - Bucket sort: input is generated by a random process that distributes elements uniformly and independently over the interval  $[0, 1)$ .

# Bucket Sort: Strategy

- Bucket sort divides the interval  $[0, 1)$  into  $n$  equal-sized subintervals, or [buckets](#).
- Then, it distributes the  $n$  input numbers into the buckets.
- Since the inputs are uniformly and independently distributed over  $[0, 1)$ , we do not expect many numbers to fall into each bucket.
- To produce the output, we simply sort the numbers in each bucket and then go through the buckets in order, listing the elements in each.

# Bucket Sort: Algorithm


- The BUCKET-SORT procedure assumes that the input is an array  $A[1:n]$  and that each element  $A[i]$  in the array satisfies  $0 \leq A[i] < 1$ .
- The code requires an auxiliary array  $B[0:n-1]$  of linked lists (buckets) and assumes that there is a mechanism for maintaining such lists.

BUCKET-SORT( $A, n$ )

```
1  let  $B[0:n-1]$  be a new array
2  for  $i = 0$  to  $n-1$ 
3      make  $B[i]$  an empty list
4  for  $i = 1$  to  $n$ 
5      insert  $A[i]$  into list  $B[\lfloor n \cdot A[i] \rfloor]$ 
6  for  $i = 0$  to  $n-1$ 
7      sort list  $B[i]$  with insertion sort
8  concatenate the lists  $B[0], B[1], \dots, B[n-1]$  together in order
9  return the concatenated lists
```

# Bucket Sort: Example

- The operation of bucket sort on an input array of 10 numbers.



# Bucket Sort: Running Time

- In the BUCKET-SORT procedure, all lines except line 7 take  $O(n)$  time in the worst case. We need to analyze the total time taken by the  $n$  calls to insertion sort in line 7.
- Let  $n_i$  be the random variable denoting the **number of elements placed in bucket  $B[i]$** . Since insertion sort runs in quadratic time, the running time of bucket sort is:

$$T(n) = \Theta(n) + \sum_{i=0}^{n-1} O(n_i^2)$$

- We now analyze the average-case running time of bucket sort by computing the expected value of the running time.

$$\begin{aligned} E[T(n)] &= E\left[\Theta(n) + \sum_{i=0}^{n-1} O(n_i^2)\right] \\ &= \Theta(n) + \sum_{i=0}^{n-1} E[O(n_i^2)] \\ &= \Theta(n) + \sum_{i=0}^{n-1} O(E[n_i^2]) \end{aligned}$$



# Bucket Sort: Running Time

- What is  $E[n_i^2]$ ?
- First of all, we can see that each bucket  $i$  has the same value of  $E[n_i^2]$ , since each value in the input array  $A$  is equally likely to fall in any bucket.
- Let's view each random variable  $n_i$  as the number of successes in  $n$  Bernoulli trials. Success in a trial occurs when an element goes into bucket  $B[i]$ .
  - Success probability  $p = 1/n$ .
  - Failure probability  $q = 1 - 1/n$ .
- A binomial distribution counts  $n_i$ , the number of successes, in the  $n$  trials.
- $E[n_i] = np$ ,  $\text{Var}[n_i] = npq$ .
- $E[n_i] = np = n \left(\frac{1}{n}\right) = 1$
- $\text{Var}[n_i] = npq = 1 - 1/n$

# Bucket Sort: Running Time

- $E[n_i^2] = \text{Var}[n_i] + E^2[n_i]$   
 $= \left(1 - \frac{1}{n}\right) + 1^2$   
 $= 2 - \frac{1}{n}$
- Thus,  $E[T(n)] = \Theta(n) + \sum_{i=0}^{n-1} O(E[n_i^2])$   
 $= \Theta(n) + n \cdot O\left(2 - \frac{1}{n}\right)$   
 $= \Theta(n).$

# Bucket Sort: Summary

- Even if the input is not drawn from a uniform distribution, bucket sort may still run in linear time.
  - As long as the input has the property that the sum of the squares of the bucket sizes is linear in the total number of elements

# End of Class

## Questions?

Instructor office: AS-1013

Email: [jso1@sogang.ac.kr](mailto:jso1@sogang.ac.kr)