

CSE3081 Design and Analysis of Algorithms

Dept. of Computer Engineering,
Sogang University

This material contains text and figures from other lecture slides. Do not post it on the Internet.

Chapter 12. Binary Search Trees

Binary Search Tree: Introduction

- A search tree data structure supports each of the following dynamic set operations

SEARCH(S, k)

A query that, given a set S and a key value k , returns a pointer x to an element in S such that $x.key = k$, or NIL if no such element belongs to S .

INSERT(S, x)

A modifying operation that adds the element pointed to by x to the set S . We usually assume that any attributes in element x needed by the set implementation have already been initialized.

DELETE(S, x)

A modifying operation that, given a pointer x to an element in the set S , removes x from S . (Note that this operation takes a pointer to an element x , not a key value.)

MINIMUM(S) and MAXIMUM(S)

Queries on a totally ordered set S that return a pointer to the element of S with the smallest (for MINIMUM) or largest (for MAXIMUM) key.

SUCCESSOR(S, x)

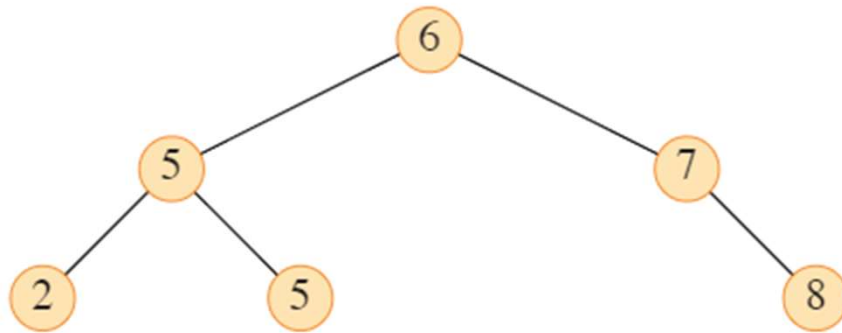
A query that, given an element x whose key is from a totally ordered set S , returns a pointer to the next larger element in S , or NIL if x is the maximum element.

PREDECESSOR(S, x)

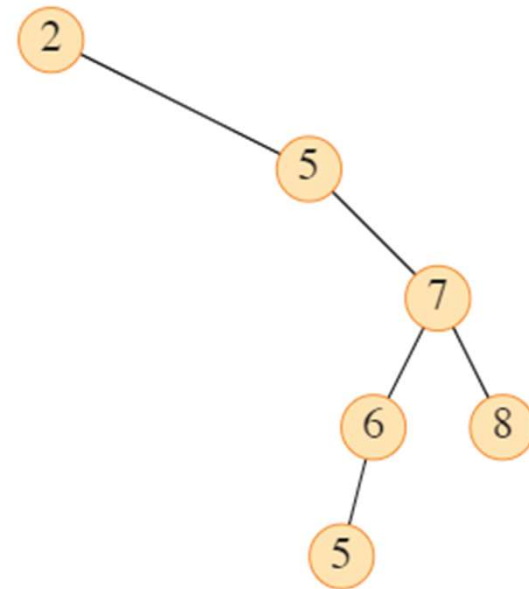
A query that, given an element x whose key is from a totally ordered set S , returns a pointer to the next smaller element in S , or NIL if x is the minimum element.

Binary Search Tree: Introduction

- Basic operations on a binary search tree take time proportional to the **height of the tree**.
 - The **height of a node** in a tree is the number of edges on the longest simple downward path from the node to a leaf.
 - The **height of a tree** is the height of its root.



height: 2



height: 4

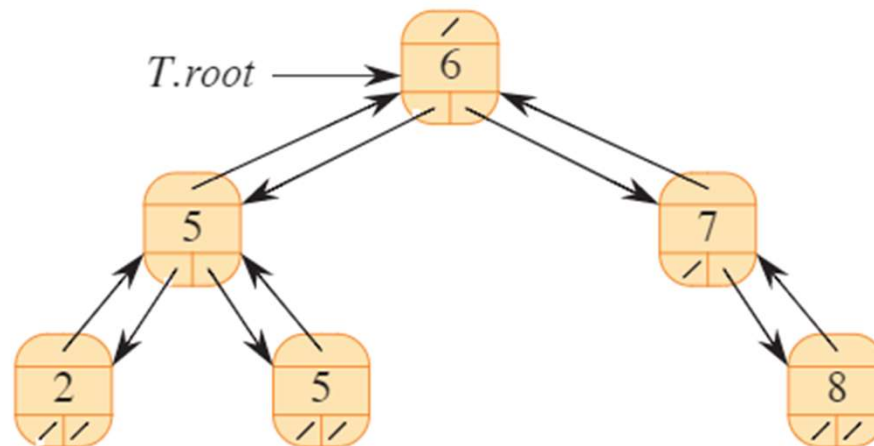
Binary Search Tree: Introduction

- For a complete binary tree with n nodes, such operations run in $\Theta(\log n)$ worst-case time.
 - A **complete binary tree** is a binary tree in which all the levels are completely filled except possibly the lowest one, which is filled from the left.
- If a tree is a linear chain of nodes, however, the same operations take $\Theta(n)$ worst-case time.

12.1 What is a binary search tree?

Binary Tree Representation

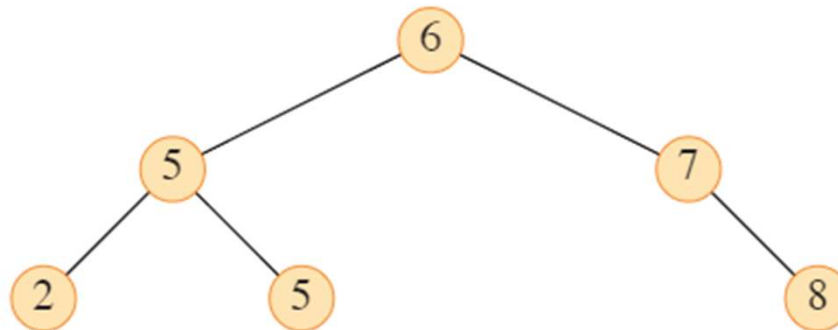
- Representing a binary tree
 - A binary tree can be represented using a linked data structure
 - Each node of a tree has the following attributes
 - *key*
 - *left*: link to left child
 - *right*: link to right child
 - *p*: link to parent
 - If a child or the parent is missing, the appropriate attribute contains NIL.
 - The tree itself has an attribute *root* that points to the root node.



Binary Search Tree: Property

- A binary search tree is a binary tree that satisfies the [binary-search-tree property](#).

Let x be a node in a binary search tree. If y is a node in the left subtree of x , then $y.key \leq x.key$. If y is a node in the right subtree of x , then $y.key \geq x.key$.



Binary Search Tree: Traversal

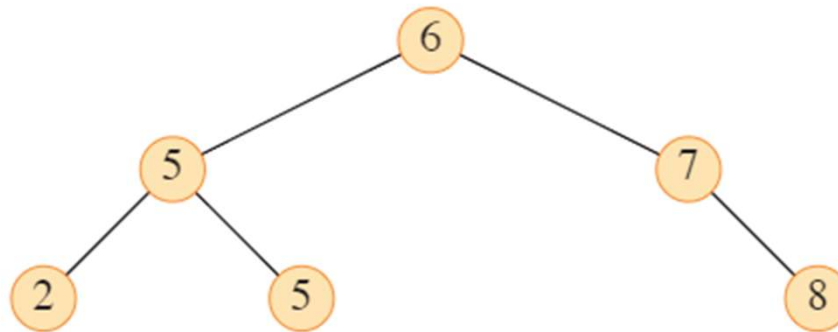
- In order to print all the keys in a sorted order, we can do the **in-order tree walk** on the tree.
- In-order tree walk is a recursive algorithm where it prints the key of the root of a subtree between printing the values in its left subtree and printing those in its right subtree.

```
INORDER-TREE-WALK( $x$ )  
1  if  $x \neq \text{NIL}$   
2      INORDER-TREE-WALK( $x.\text{left}$ )  
3      print  $x.\text{key}$   
4      INORDER-TREE-WALK( $x.\text{right}$ )
```

- Other walk algorithms
 - pre-order walk: print root before the values in either subtree
 - post-order walk: print root after the values in either subtree

Binary Search Tree: Example Tree Walk

- In-order tree walk on a binary search tree
 - $2 \rightarrow 5 \rightarrow 5 \rightarrow 6 \rightarrow 7 \rightarrow 8$
- Pre-order walk
 - $6 \rightarrow 5 \rightarrow 2 \rightarrow 5 \rightarrow 7 \rightarrow 8$
- Post-order walk
 - $2 \rightarrow 5 \rightarrow 5 \rightarrow 8 \rightarrow 7 \rightarrow 6$



Binary Search Tree: Cost of Walk

- It takes $\Theta(n)$ time to walk an n -node binary search tree
- After the initial call, the procedure calls itself recursively exactly twice for each node in the tree – once for its left child and once for its right child.
- Formal proof that the walk takes $O(n)$ time.
 - Suppose that INORDER-TREE-WALK is called on a node x whose left subtree has k nodes and whose right subtree has $n - k - 1$ nodes.
 - $T(n) \leq T(k) + T(n - k - 1) + d$
 - d is a constant that reflects an upper bound on the time to execute the body of INORDER-TREE-WALK, exclusive of the time spent in recursive calls.
 - Using substitution method, we show that $T(n) \leq (c + d)n + c$. ($T(0) = c$)

$$\begin{aligned} T(n) &\leq T(k) + T(n - k - 1) + d \\ &\leq ((c + d)k + c) + ((c + d)(n - k - 1) + c) + d \\ &= (c + d)n + c - (c + d) + c + d \\ &= (c + d)n + c, \end{aligned}$$

12.2 Querying a binary search tree

Binary Search Tree: Querying Operations

- Binary search trees can support the queries MINIMUM, MAXIMUM, SUCCESSOR, PREDECESSOR, as well as SEARCH.
- These operations can be supported in $O(h)$ time where h is the height of the tree.

Binary Search Tree: Searching

- Given a pointer x to the root of a subtree and a key k
- $\text{TREE-SEARCH}(x, k)$ returns a pointer to a node with key k if one exists in the subtree; otherwise, it returns NIL.
- Recursive implementation

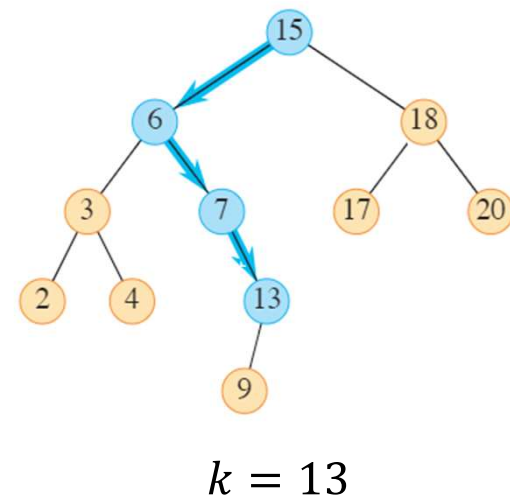
```
TREE-SEARCH( $x, k$ )  
1  if  $x == \text{NIL}$  or  $k == x.\text{key}$   
2      return  $x$   
3  if  $k < x.\text{key}$   
4      return  $\text{TREE-SEARCH}(x.\text{left}, k)$   
5  else return  $\text{TREE-SEARCH}(x.\text{right}, k)$ 
```

Binary Search Tree: Searching

- Given a pointer x to the root of a subtree and a key k
- $\text{TREE-SEARCH}(x, k)$ returns a pointer to a node with key k if one exists in the subtree; otherwise, it returns NIL.
- Recursive implementation

$\text{TREE-SEARCH}(x, k)$

```
1  if  $x == \text{NIL}$  or  $k == x.\text{key}$ 
2      return  $x$ 
3  if  $k < x.\text{key}$ 
4      return  $\text{TREE-SEARCH}(x.\text{left}, k)$ 
5  else return  $\text{TREE-SEARCH}(x.\text{right}, k)$ 
```



- The nodes encountered during the recursion form a simple path downward from the root of the tree → **running time of TREE-SEARCH is $O(h)$.**

Binary Search Tree: Searching

- Iterative implementation

```
ITERATIVE-TREE-SEARCH( $x, k$ )  
1  while  $x \neq \text{NIL}$  and  $k \neq x.\text{key}$   
2      if  $k < x.\text{key}$   
3           $x = x.\text{left}$   
4      else  $x = x.\text{right}$   
5  return  $x$ 
```


Binary Search Tree: Minimum and Maximum

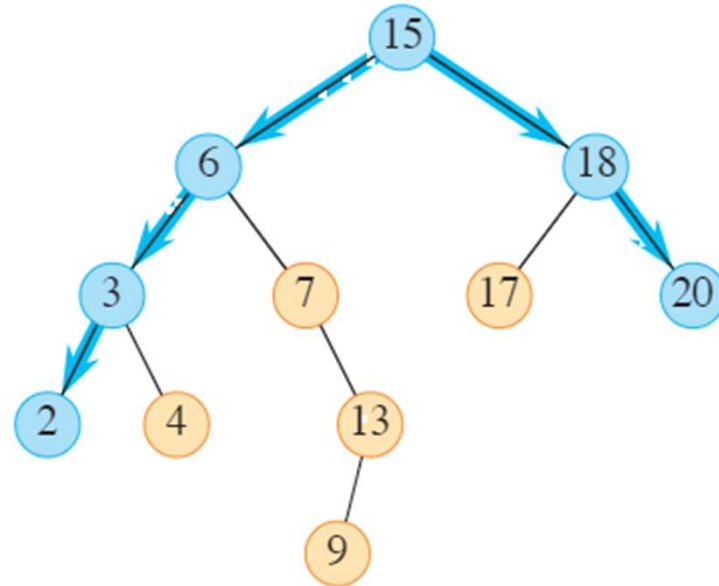
- To find the minimum key in the tree, just follow the left child pointers from the root until you encounter a NIL.
- To find the maximum key in the tree, just follow the right child pointers from the root until you encounter a NIL.

TREE-MINIMUM(x)

```
1  while  $x.left \neq \text{NIL}$ 
2       $x = x.left$ 
3  return  $x$ 
```

TREE-MAXIMUM(x)

```
1  while  $x.right \neq \text{NIL}$ 
2       $x = x.right$ 
3  return  $x$ 
```



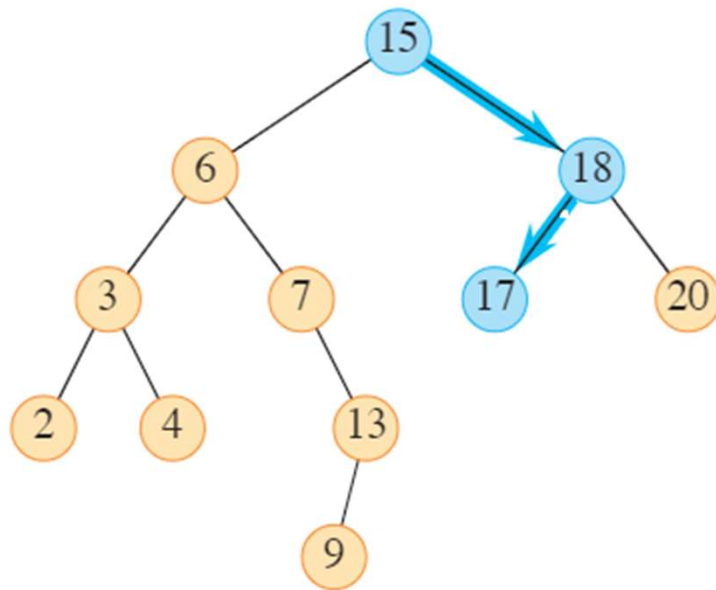
- Similar to searching, finding minimum or maximum takes $\Theta(h)$ time.

Binary Search Tree: Successor and Predecessor

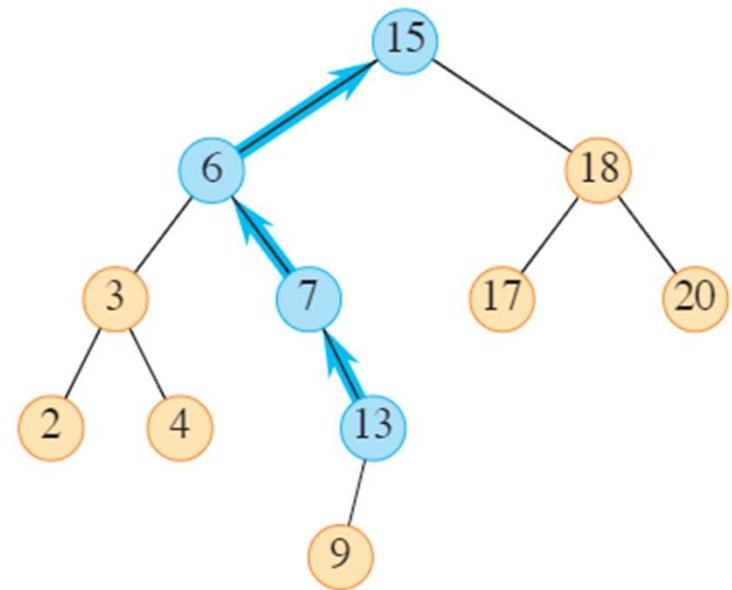
- If all keys are distinct, the successor of a node x is the node with the smallest key greater than $x.key$.
- In a binary search tree, the **successor** of a node is the **next node visited in an in-order tree walk**.
- There are 2 case when finding successor of node x in a binary search tree.
- Case 1: If the right subtree of node x is nonempty, then the successor of x is the **leftmost node in x 's right subtree**.
- Case 2: if the right subtree of node x is empty and x has a successor y , then **y is the lowest ancestor of x whose left child is also an ancestor of x** .

Binary Search Tree: Successor and Predecessor

- Example



right child is nonempty



right child is empty

Binary Search Tree: Successor and Predecessor

- TREE-SUCCESSOR procedure

```
TREE-SUCCESSOR( $x$ )
1  if  $x.right \neq \text{NIL}$ 
2      return TREE-MINIMUM( $x.right$ )  // leftmost node in right subtree
3  else // find the lowest ancestor of  $x$  whose left child is an ancestor of  $x$ 
4       $y = x.p$ 
5      while  $y \neq \text{NIL}$  and  $x == y.right$ 
6           $x = y$ 
7           $y = y.p$ 
8      return  $y$ 
```

- The running time of TREE-SUCCESSOR is $O(h)$, since it either follows a simple path up the tree or follows a simple path down the tree.
- TREE-PREDECESSOR is symmetric to TREE-SUCCESSOR.

12.3 Insertion and deletion

Insertion and Deletion

- When we insert an element to the binary search tree or delete an element from the tree, we modify the binary search tree structure.
- Insertion and deletion must be done in a way that the **binary-search-tree property continues to hold** after insertion or deletion.

Insertion

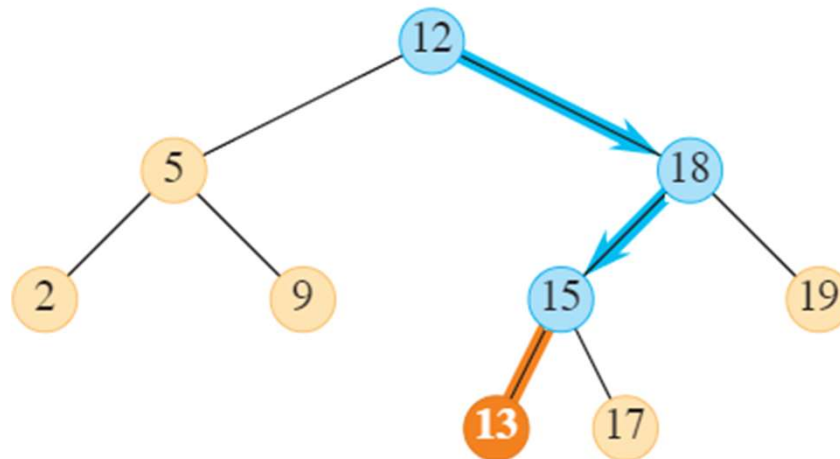
- The TREE-INSERT procedure inserts a new node into a binary search tree.
- The procedure takes a binary tree T and a node z for which $z.key$ has already been filled, $z.left = \text{NIL}$ and $z.right = \text{NIL}$.
- z must be inserted into an appropriate position in the tree.

TREE-INSERT(T, z)

```
1   $x = T.root$            // node being compared with  $z$ 
2   $y = \text{NIL}$              //  $y$  will be parent of  $z$ 
3  while  $x \neq \text{NIL}$       // descend until reaching a leaf
4       $y = x$ 
5      if  $z.key < x.key$ 
6           $x = x.left$ 
7      else  $x = x.right$ 
8   $z.p = y$                // found the location—insert  $z$  with parent  $y$ 
9  if  $y == \text{NIL}$ 
10      $T.root = z$         // tree  $T$  was empty
11 elseif  $z.key < y.key$ 
12      $y.left = z$ 
13 else  $y.right = z$ 
```

Insertion

- The procedure maintains a pointer x that traces a simple path downward looking for a NIL to replace with the input node z .
- The procedure also maintains a trailing pointer y as the parent of x .
- Once x finds NIL, z should be inserted in that position.
- x is inserted into the tree as y 's child.



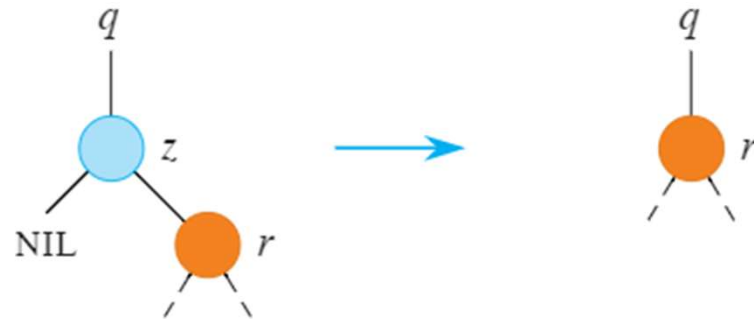
- Like other operations, it is easy to see that TREE-INSERT runs in $O(h)$ time.

Deletion

- The overall strategy for deleting a node z from a binary search tree T has three basic cases.
- Case 1: If z has no children, then simply remove it by modifying its parent to replace z with NIL as its child.
- Case 2: If z has just one child, then elevate that child to take z 's position in the tree by modifying z 's parent to replace z by z 's child.
- Case 3: If z has two children, find z 's successor y – which must belong to z 's right subtree – and move y to take z 's position in the tree.
 - The rest of z 's original right subtree becomes y 's new right subtree, and z 's left subtree becomes y 's new left subtree.
 - Because y is z 's successor, it cannot have a left child, and y 's original right child moves into y 's original position, with the rest of y 's original right subtree following automatically.

Deletion: Case 2

- If z has no left child, replace z by its right child.
 - If right child is NIL, it is case 1.

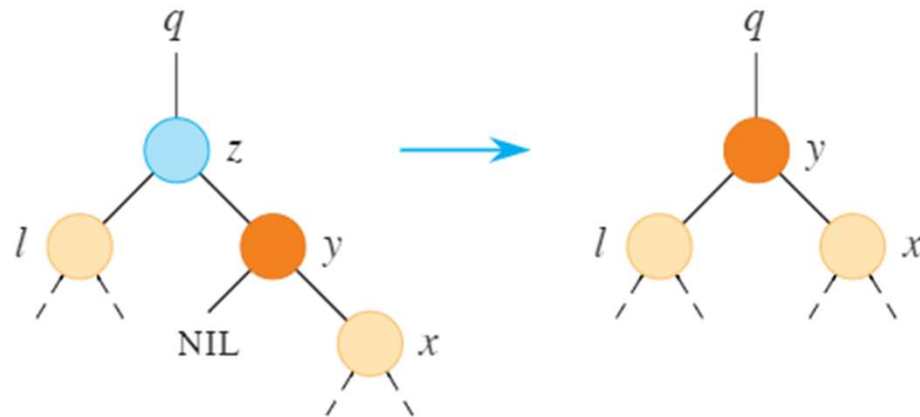


- If z has no right child, place z by its left child.
 - If left child is NIL, it is case 1.



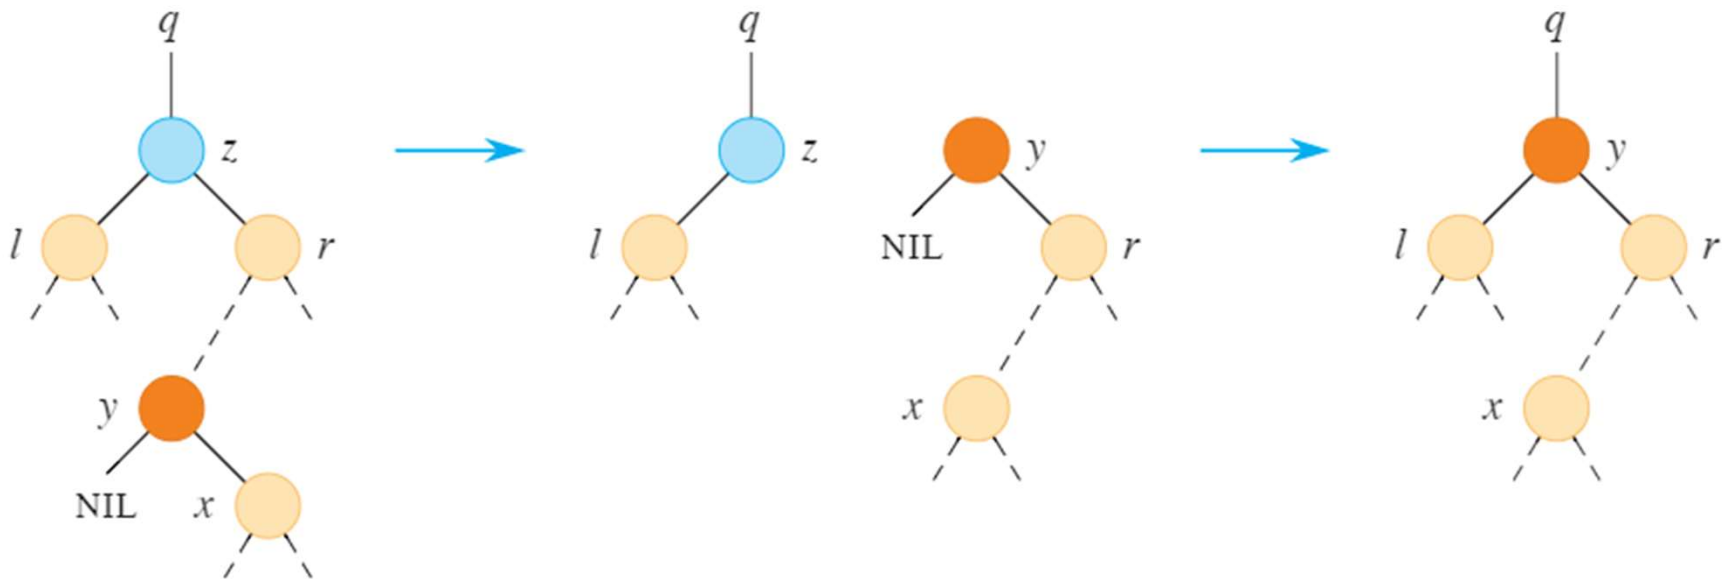
Deletion: Case 3-i

- If the successor y is z 's right child, replace z by y , leaving y 's child alone.



Deletion: Case 3-ii

- Otherwise, y lies within z 's right subtree but is not z 's right child. In this case, first replace y by its own right child, and then replace z by y .



Deletion: TRANSPLANT

- The subroutine TRANSPLANT replaces the subtree rooted at node u with the subtree rooted at node v , node u 's parent becomes node v 's parent and u 's parent ends up having v as its appropriate child.
 - v can be NIL.

TRANSPLANT(T, u, v)

```
1  if  $u.p == \text{NIL}$ 
2       $T.root = v$ 
3  elseif  $u == u.p.left$ 
4       $u.p.left = v$ 
5  else  $u.p.right = v$ 
6  if  $v \neq \text{NIL}$ 
7       $v.p = u.p$ 
```

Deletion: TREE-DELETE

- TREE-DELETE removes node z from the binary search tree T .

```
TREE-DELETE( $T, z$ )
1  if  $z.left == \text{NIL}$ 
2      TRANSPLANT( $T, z, z.right$ )           // replace  $z$  by its right child
3  elseif  $z.right == \text{NIL}$ 
4      TRANSPLANT( $T, z, z.left$ )             // replace  $z$  by its left child
5  else  $y = \text{TREE-MINIMUM}(z.right)$         //  $y$  is  $z$ 's successor
6      if  $y \neq z.right$                     // is  $y$  farther down the tree?
7          TRANSPLANT( $T, y, y.right$ )        // replace  $y$  by its right child
8           $y.right = z.right$                 //  $z$ 's right child becomes
9           $y.right.p = y$                     //  $y$ 's right child
10     TRANSPLANT( $T, z, y$ )                  // replace  $z$  by its successor  $y$ 
11      $y.left = z.left$                       // and give  $z$ 's left child to  $y$ ,
12      $y.left.p = y$                         // which had no left child
```

- All lines except for the call to TREE-MINIMUM takes constant time.
- Thus, TREE-DELETE runs in $O(h)$ time.

End of Class

Questions?

Instructor office: AS-1013

Email: jso1@sogang.ac.kr